

Even more hierarchical models

FW 891

[Click here to view presentation online](#)

Christopher Cahill

25 October 2023



Quantitative Fisheries Center
MICHIGAN STATE UNIVERSITY

Purpose

- Today we introduce a powerful extension of mixed effects models
 - Random slopes (aka varying effects models)
 - Repent for our earlier sins
- Adventures in covariance a la McElreath (2023)
 - Cover some math necessary for working with covariance matrices
- Simulate a varying effects problem
- Develop both centered and noncentered varying effects models in Stan

Some references

- This lecture is drawing heavily on McElreath (2023), and much of the code and analyses are adapted from information in that text

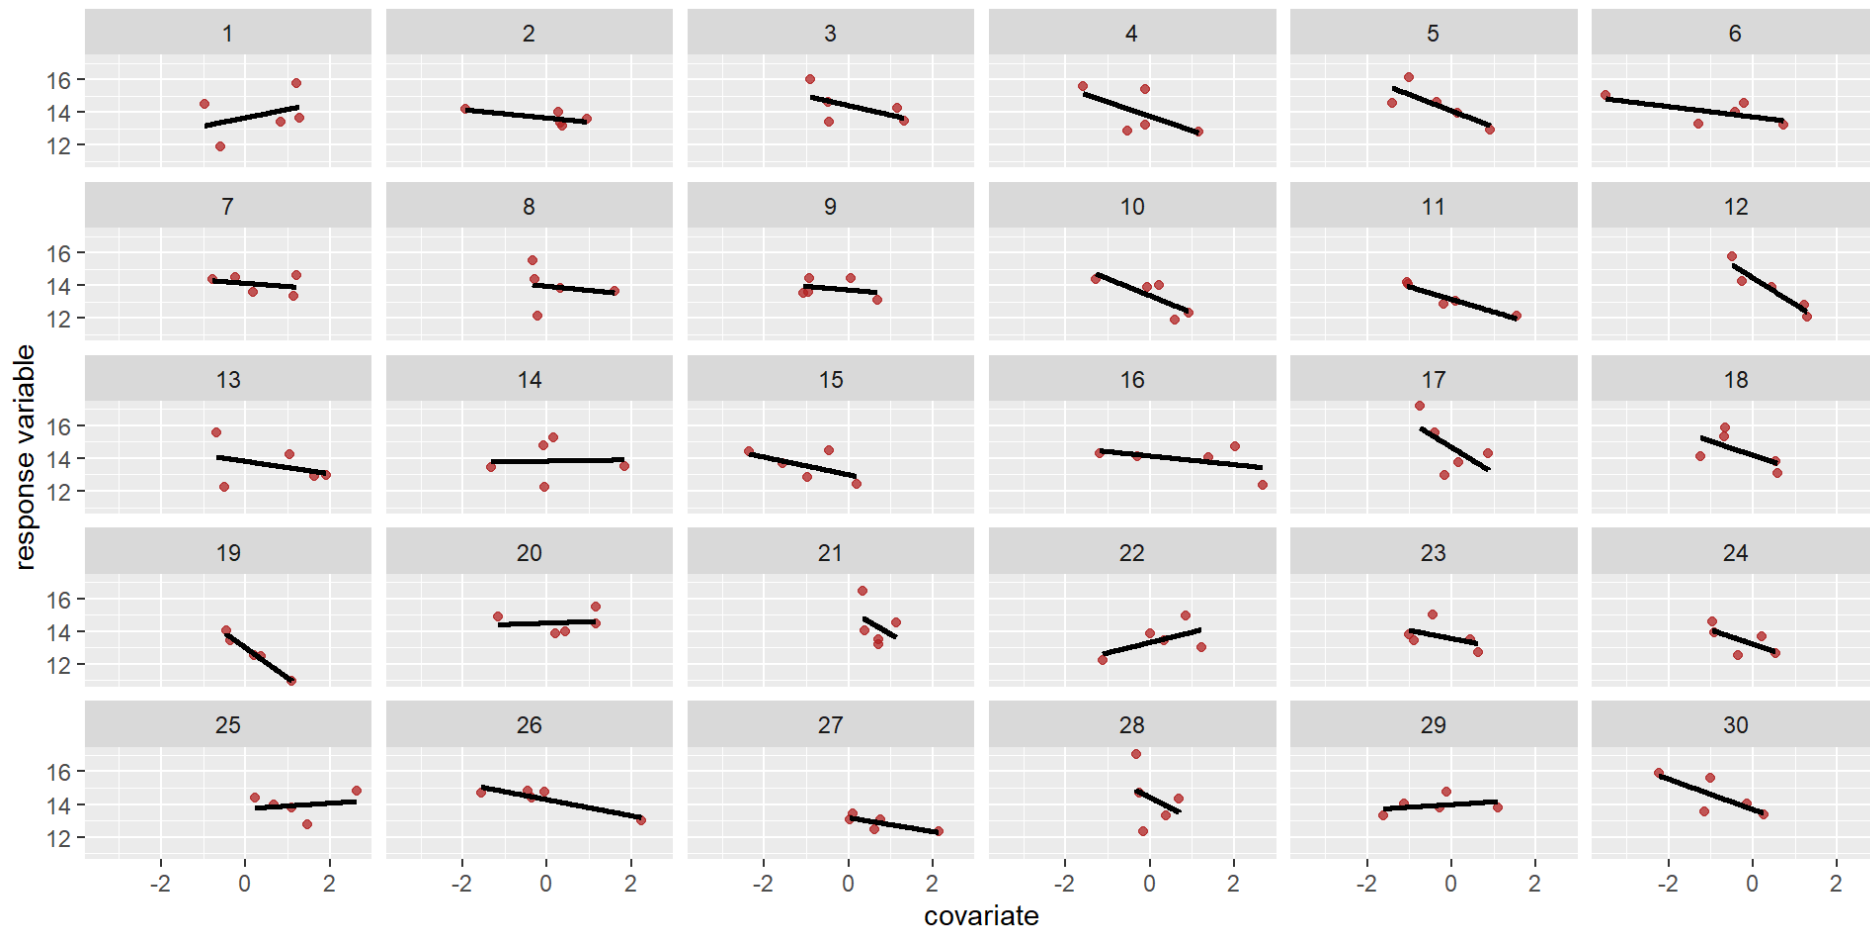
Some references

- This lecture is drawing heavily on McElreath (2023), and much of the code and analyses are adapted from information in that text
 - If you want more information, check “Adventures in Covariance” chapter of this book

Some references

- This lecture is drawing heavily on McElreath (2023), and much of the code and analyses are adapted from information in that text
 - If you want more information, check “Adventures in Covariance” chapter of this book
- See also Gelman and Hill (2007)
 - Specifically chapter 13

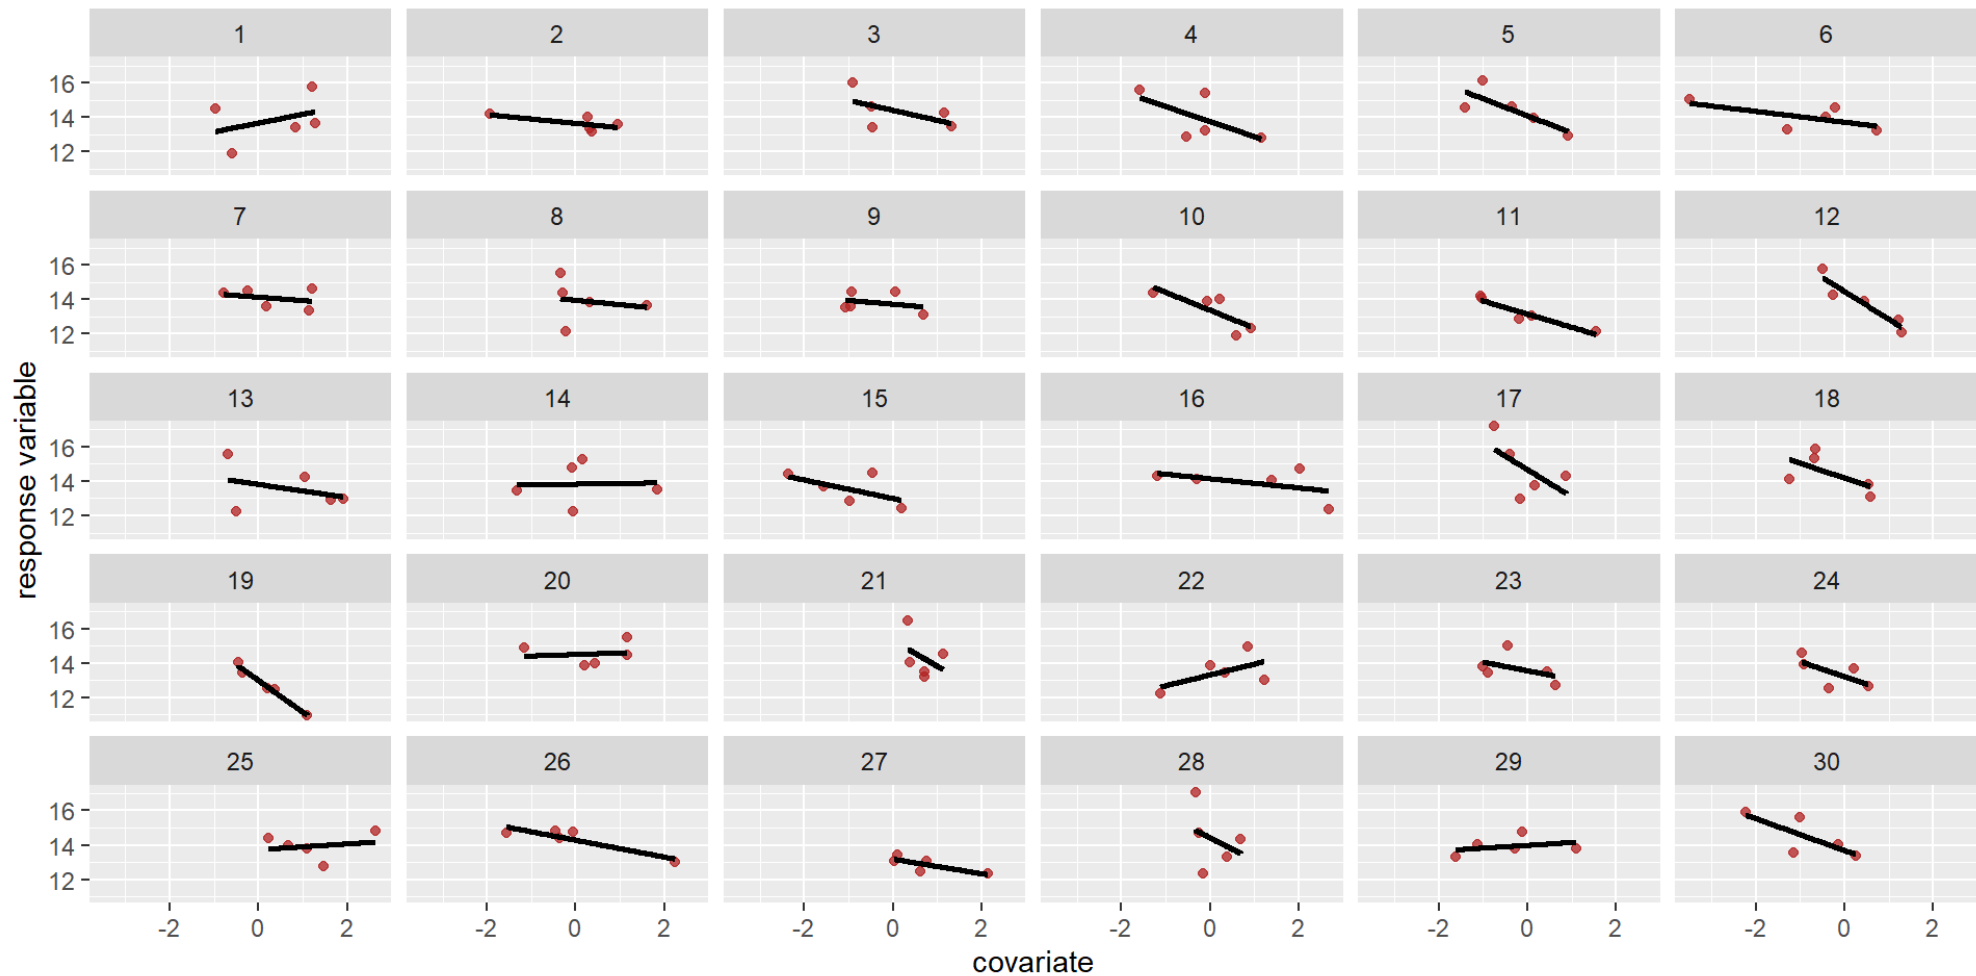
Thinking about variability in ecological systems



Three key takeaways


Point #1

- variability in both intercepts *and* slopes among replicates




Point #2

- slopes get steeper as intercepts get bigger



Point #3


- some groups display **Simpson's paradox**



Simpson's paradox

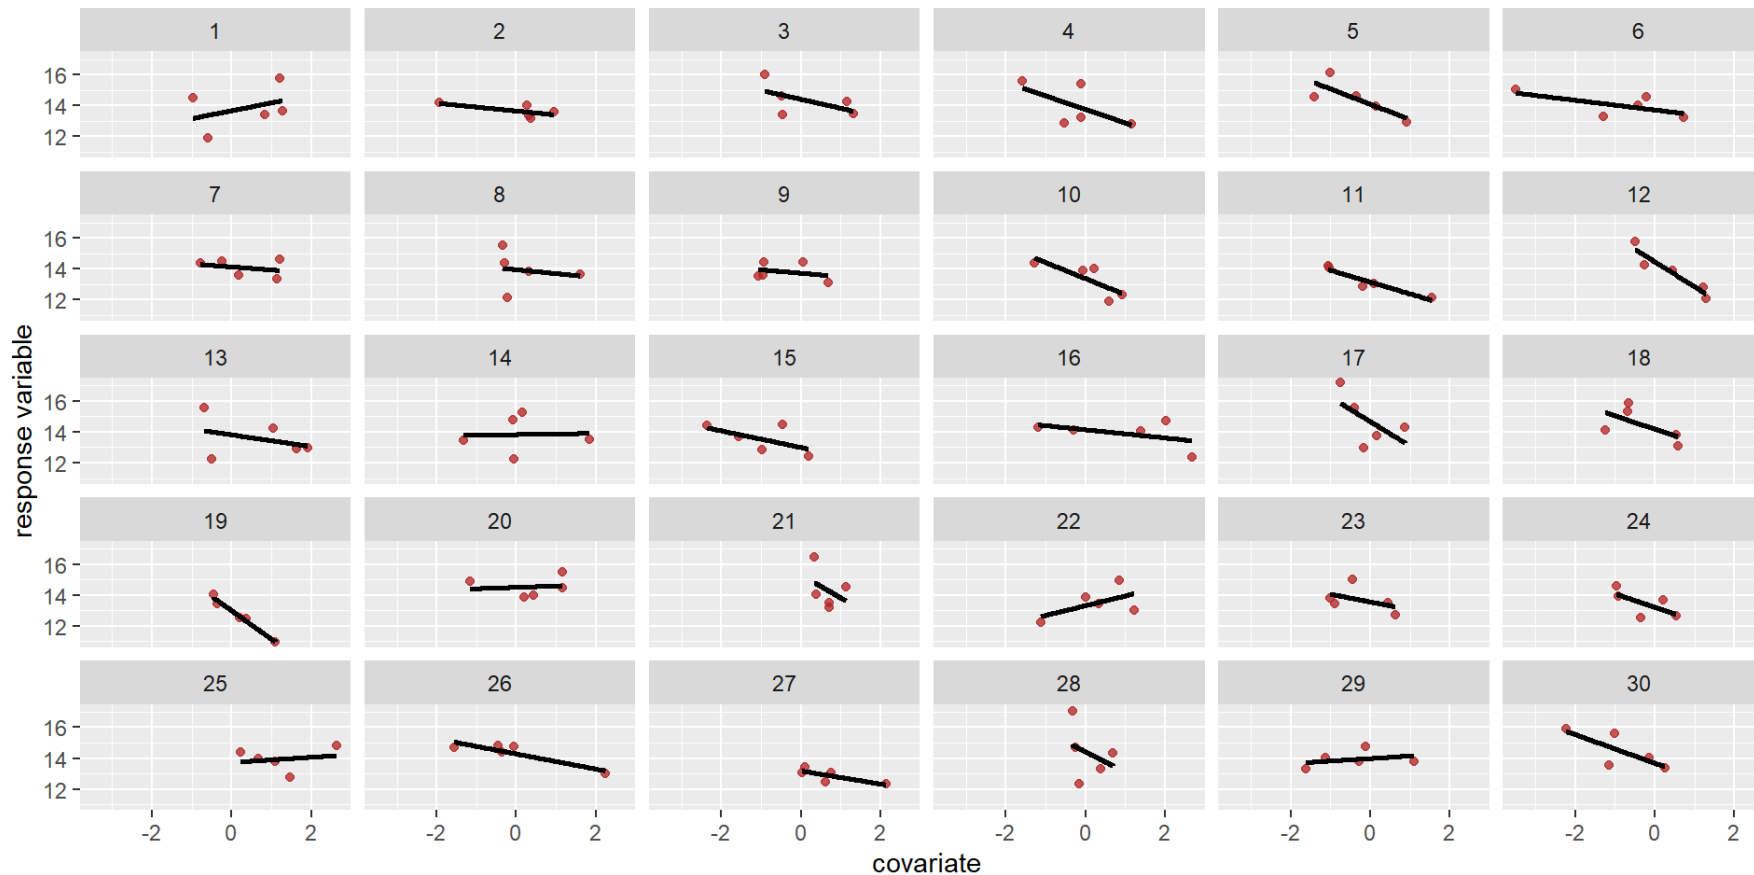
Colored points are data from replicate 1

Grey points are data from all other replicates



Group question

- what is a question in your field of study that might show a similar pattern?



Varying effects

- Generalization of standard multilevel regression
 - Specifically, models that allow slopes and intercepts to vary by group
- Several ways to write, here's one for the model we just visualized

Varying effects maths

$$y_i \sim \text{N}(\mu_i, \sigma) \quad [\text{likelihood}]$$

$$\mu_i = \beta_{0[\textit{group}]} + \beta_{1[\textit{group}]} x_{1[i]} \quad [\text{linear model}]$$

Varying effects maths

$$y_i \sim \text{N}(\mu_i, \sigma) \quad [\text{likelihood}]$$

$$\mu_i = \beta_{0[\textit{group}]} + \beta_{1[\textit{group}]} x_{1[i]} \quad [\text{linear model}]$$

Then comes the matrix of varying intercepts and slopes, with its covariance matrix Σ :

Varying effects maths

$$y_i \sim \mathcal{N}(\mu_i, \sigma) \quad [\text{likelihood}]$$

$$\mu_i = \beta_{0[group]} + \beta_{1[group]} x_{1[i]} \quad [\text{linear model}]$$

Then comes the matrix of varying intercepts and slopes, with its covariance matrix Σ :

$$\begin{bmatrix} \beta_{0_{group}} \\ \beta_{1_{group}} \end{bmatrix} \sim \text{MVN} \left(\begin{bmatrix} \beta_0 \\ \beta_1 \end{bmatrix}, \Sigma \right) \quad [\text{population of varying effects}]$$

$$\Sigma = \begin{pmatrix} \sigma_{\beta_0} & 0 \\ 0 & \sigma_{\beta_1} \end{pmatrix} \Omega \begin{pmatrix} \sigma_{\beta_0} & 0 \\ 0 & \sigma_{\beta_1} \end{pmatrix} \quad [\text{construct covariance matrix}]$$

Varying effects maths

$$y_i \sim \mathcal{N}(\mu_i, \sigma) \quad [\text{likelihood}]$$

$$\mu_i = \beta_{0[group]} + \beta_{1[group]} x_{1[i]} \quad [\text{linear model}]$$

Then comes the matrix of varying intercepts and slopes, with its covariance matrix Σ :

$$\begin{bmatrix} \beta_{0_{group}} \\ \beta_{1_{group}} \end{bmatrix} \sim \text{MVN} \left(\begin{bmatrix} \beta_0 \\ \beta_1 \end{bmatrix}, \Sigma \right) \quad [\text{population of varying effects}]$$

$$\Sigma = \begin{pmatrix} \sigma_{\beta_0} & 0 \\ 0 & \sigma_{\beta_1} \end{pmatrix} \Omega \begin{pmatrix} \sigma_{\beta_0} & 0 \\ 0 & \sigma_{\beta_1} \end{pmatrix} \quad [\text{construct covariance matrix}]$$

- Ω is a correlation matrix

Let's break that down

- Math on the previous slide says that each group has a $\beta_{0[group]}$ and $\beta_{1[group]}$ with a prior distribution defined by the two dimensional Gaussian distribution with means β_0 and β_1 and covariance matrix Σ

Let's break that down

- Math on the previous slide says that each group has a $\beta_{0[group]}$ and $\beta_{1[group]}$ with a prior distribution defined by the two dimensional Gaussian distribution with means β_0 and β_1 and covariance matrix Σ
 - This is a multivariate normal distribution

Let's break that down

- Math on the previous slide says that each group has a $\beta_{0[group]}$ and $\beta_{1[group]}$ with a prior distribution defined by the two dimensional Gaussian distribution with means β_0 and β_1 and covariance matrix Σ
 - This is a multivariate normal distribution
- The final line defines how we construct our covariance matrix Σ by factoring it into a diagonal matrix of σ_{β_0} and σ_{β_1} and a correlation matrix Ω

Let's break that down

- Math on the previous slide says that each group has a $\beta_{0[group]}$ and $\beta_{1[group]}$ with a prior distribution defined by the two dimensional Gaussian distribution with means β_0 and β_1 and covariance matrix Σ
 - This is a multivariate normal distribution
- The final line defines how we construct our covariance matrix Σ by factoring it into a diagonal matrix of σ_{β_0} and σ_{β_1} and a correlation matrix Ω
 - Several ways to construct Σ , but splitting it into standard deviations, σ_{β_0} and σ_{β_1} , and a correlation matrix Ω helps with understanding

Let's break that down

- Math on the previous slide says that each group has a $\beta_{0[group]}$ and $\beta_{1[group]}$ with a prior distribution defined by the two dimensional Gaussian distribution with means β_0 and β_1 and covariance matrix Σ
 - This is a multivariate normal distribution
- The final line defines how we construct our covariance matrix Σ by factoring it into a diagonal matrix of σ_{β_0} and σ_{β_1} and a correlation matrix Ω
 - Several ways to construct Σ , but splitting it into standard deviations, σ_{β_0} and σ_{β_1} , and a correlation matrix Ω helps with learning
- Compare this with a standard normal distribution which takes a mean and a standard deviation

The correlation matrix

- For this simple example, the correlation matrix looks like

$$\Omega = \begin{pmatrix} 1 & \rho \\ \rho & 1 \end{pmatrix}$$

The correlation matrix

- For this simple example, the correlation matrix looks like

$$\mathbf{\Omega} = \begin{pmatrix} 1 & \rho \\ \rho & 1 \end{pmatrix}$$

- where ρ is the correlation between β_0 and β_1

The correlation matrix

- For this simple example, the correlation matrix looks like

$$\Omega = \begin{pmatrix} 1 & \rho \\ \rho & 1 \end{pmatrix}$$

- where ρ is the correlation between β_0 and β_1
- Ω gets more complicated for models with more varying slopes

Cholesky decomposition

- Note that we can take any arbitrary symmetric, positive-definite matrix \mathbf{A} , and factor or decompose it into

$$\mathbf{A} = \mathbf{L}\mathbf{L}^T$$

Cholesky decomposition

- Note that we can take any arbitrary symmetric, positive-definite matrix \mathbf{A} , and factor or decompose it into

$$\mathbf{A} = \mathbf{L}\mathbf{L}^T$$

- where \mathbf{L} is a lower triangular matrix with real and positive diagonal entries and \mathbf{L}^T is a transpose of \mathbf{L}

Cholesky Decomposition

- If we visualize a Cholesky decomposition

$$\begin{bmatrix} A_{00} & A_{01} & A_{02} \\ A_{10} & A_{11} & A_{12} \\ A_{20} & A_{21} & A_{22} \end{bmatrix} = \begin{bmatrix} L_{00} & 0 & 0 \\ L_{10} & L_{11} & 0 \\ L_{20} & L_{21} & L_{22} \end{bmatrix} \begin{bmatrix} L_{00} & L_{10} & L_{20} \\ 0 & L_{11} & L_{21} \\ 0 & 0 & L_{22} \end{bmatrix}$$

Cholesky Decomposition

- If we visualize a Cholesky decomposition

$$\begin{bmatrix} A_{00} & A_{01} & A_{02} \\ A_{10} & A_{11} & A_{12} \\ A_{20} & A_{21} & A_{22} \end{bmatrix} = \begin{bmatrix} L_{00} & 0 & 0 \\ L_{10} & L_{11} & 0 \\ L_{20} & L_{21} & L_{22} \end{bmatrix} \begin{bmatrix} L_{00} & L_{10} & L_{20} \\ 0 & L_{11} & L_{21} \\ 0 & 0 & L_{22} \end{bmatrix}$$

- This is helpful from a numerical perspective, particularly with noncentered parameterizations

Cholesky factors continued

- Note that there is a lot of convenient linear algebra that can be done with Cholesky factors of covariance matrices \mathbf{L} or of correlation matrices \mathbf{L}_{corr}
 - For example,

$$\mathbf{L} = \begin{pmatrix} \sigma_{\beta_0} & 0 \\ 0 & \sigma_{\beta_1} \end{pmatrix} \mathbf{L}_{\text{corr}}$$

- See this [link](#) for a useful review

Cholesky factors

```
1 # create a correlation matrix and declare sigmas
2 OMEGA <- matrix(c(1, 0.7, 0.7, 1), nrow = 2)
3 sigmas <- c(1, 2) # sd_b0, sd_b1
4
5 OMEGA
```

```
      [,1] [,2]
[1,]  1.0  0.7
[2,]  0.7  1.0
```

```
1 sigmas
```

```
[1] 1 2
```

```
1 # note also
2 diag(sigmas) # diagonal matrix
```

```
      [,1] [,2]
[1,]    1    0
[2,]    0    2
```

Cholesky factors

```
1 # calculate covariance matrix:  
2 SIGMA <- diag(sigmas) %*% OMEGA %*% diag(sigmas)  
3 SIGMA
```

```
      [,1] [,2]  
[1,]  1.0  1.4  
[2,]  1.4  4.0
```


Cholesky factors

```
1 # convert Cholesky factor of correlation matrix to covariance Cholesky fact
2 L_corr <- t(chol(OMEGA)) # note chol() returns upper triangular matrix
3 diag(sigmas) %*% L_corr
```

```
      [,1]      [,2]
[1,]  1.0  0.000000
[2,]  1.4  1.428286
```

```
1 t(chol(SIGMA)) # L of SIGMA
```

```
      [,1]      [,2]
[1,]  1.0  0.000000
[2,]  1.4  1.428286
```

Cholesky factors

```
1 # Cholesky factor of correlation matrix to
2 # covariance matrix Cholesky factor:
3 L_corr <- t(chol(OMEGA)) # note chol() returns upper tri
4 Lambda <- diag(sigmas) %*% L_corr
5
6 t(chol(SIGMA)) # L of SIGMA
```



```
      [,1]      [,2]
[1,]  1.0 0.000000
[2,]  1.4 1.428286
```

```
1 Lambda
```

```
      [,1]      [,2]
[1,]  1.0 0.000000
[2,]  1.4 1.428286
```

Cholesky factors

```
1 # generate random values with desired covariance
2 Z <- rbind(rnorm(1e4), rnorm(1e4))
3 X <- Lambda %*% Z
4 par(mfrow=c(1,2))
5 plot(Z[,1]~Z[,2], main = "uncorrelated deviates", col = "firebrick")
6 plot(X[,1]~X[,2], main = "correlated deviates", col = "firebrick")
```



A problem

- People want to know the extent to which juvenile walleye growth rate is density dependent
 - Has implications for both basic ecology and management
- DNR Biologists go to a collection of lakes and measure length of age-0 walleye in fall as a proxy of juvenile growth rate
- Each year, the biologists attempt to go to 30 lakes in total (weather pending)
 - They also conduct surveys to get an estimate of juvenile density
- Let's simulate some fake data representing this problem, and then build some Stan models to recover
- go to the `varying_effects.r` script

Hyperpriors for varying effects model

$\beta_0 \sim \text{Normal}(0, 25)$ [prior for average intercept]

$\beta_1 \sim \text{Normal}(0, 25)$ [prior for average slope]

$\sigma \sim \text{Exponential}(0.01)$ [prior for stddev within group]

$\sigma_{\beta_0} \sim \text{Exponential}(0.01)$ [prior stddev among intercepts]


$\sigma_{\beta_1} \sim \text{Exponential}(0.01)$ [prior stddev among intercepts]

$\Omega \sim \text{LKJ corr}(2)$ [prior for correlation matrix]

- $\text{LKJ corr}(2)$ defines a weakly informative prior on ρ that is skeptical of extreme correlations near -1 or 1


Visualizing the LKJcorr prior

- LKJ corr(1)




Visualizing the LKJcorr prior

- LKJ corr(2)




Visualizing the LKJcorr prior

- LKJ corr(4)



Visualizing the LKJcorr prior

- LKJ corr(10)



Wrap up

- Introduced a powerful extension to mixed effects models: varying effects
- Went through a bunch of math to show how to play with multivariate normal distributions
- Simulated an example, conducted prior predictive checks, estimated the model
- Showed a non-centered version of this model as well
- Up next, spatial random effects

References

Cahill et al. 2020. A spatial-temporal approach to modeling somatic growth across inland fisheries landscapes. CJFAS.

Gelman, A. and J. Hill. 2007. Data analysis using regression and multilevel/hierarchical models

McElreath 2023. Statistical Rethinking.

Simpson's paradox Wikipedia:

https://en.wikipedia.org/wiki/Simpson%27s_paradox

<https://mlisi.xyz/post/simulating-correlated-variables-with-the-cholesky-factorization/>

